

# CASHEL FITZGERALD

[cashel@utexas.edu](mailto:cashel@utexas.edu)

[linkedin.com/in/cashelfitzgerald](https://linkedin.com/in/cashelfitzgerald)

[github.com/CashelF](https://github.com/CashelF)

[cashel.dev](https://cashel.dev)

## Education

Cornell University, NYC — Master of Engineering in Computer Science	2025 – 2026
University of Texas at Austin — B.S. in Electrical & Computer Engineering (GPA 3.8)	2021 – 2025
<b>Experience</b>	

Modern Intelligence	March 2025 – Present
Machine Learning Engineer Intern	Austin, TX

- Engineered a distributed training pipeline (AWS multi-node) processing 10M+ custom-labeled frames, utilizing knowledge distillation to compress state-of-the-art transformers for edge deployment.
- Fine-tuned large-scale transformer detectors (Co-DETR, RF-DETR) on custom aerial datasets, achieving 0.47 mAP and optimizing inference latency by 30% via operator fusion and BF16 quantization.
- Built a pre-label script linking our Deformable DETR model to Encord for auto-annotations, cutting labeler workload

8am (formerly AffiniPay)	May 2024 – August 2024
Software Engineering Intern	Austin, TX

Provider of advanced payment processing and project management solutions for professional services firms.	
• Key contributor to the technical system design that guides the evolution of how AffiniPay bills merchants (\$300M+ annual revenue, \$20B+ yearly transactions) using event-driven architecture.	
• Created event reactor for Java service to consume Kafka events with merchant data and update billing accounts.	
• Implemented major behavioral change to Scala microservice to handle HTTP redirects and upload external files to S3.	
• Built microservice to propagate Kafka partition and offset info, enabling async event tracing through the API gateway.	

Favor Delivery	June 2023 – August 2023
Software Engineering Intern	Austin, TX

On-demand delivery platform connecting customers to efficient delivery services from their favorite local businesses.	
• Implemented a real-time metrics tracking system using Micrometer, Prometheus, and Grafana for integrated merchants.	
• Automated error monitoring with cross-service alerts, reducing integration failure lead time by 99.3%.	
• Refactored the merchant POS system integrations logging infrastructure, enhancing efficiency and maintainability.	

Voyant	May 2022 – August 2022
Software Developer Intern	Austin, TX

Offers tools that support financial advisors' businesses, with a complete suite of investment solutions.	
--	--

## Extracurriculars and Projects

Activities: Texas Luminescence (Director of Engineering), Sanger Learning Center Calculus Tutor, IEEE @ UT	
Agentic Knowledge Graph Framework   Python, LangChain, Neo4j, LLMs	January 2026
• Architected an autonomous agent system capable of grounding LLM actions via a procedural Knowledge Graph, enabling complex multi-step desktop automation.	
• Implemented a Graph-RAG pipeline to retrieve structured context for reasoning, significantly reducing hallucination rates in few-shot imitation learning tasks.	
• Designed a shared workflow memory system, allowing agents to generalize learned tasks across different environments.	
Test-Time Speaker Adaptation in VSR   Transformers, Computer Vision, Multimodal	December 2025
• Built a hypernetwork to output per-sample, gradient-free LoRA for speaker adaptation in visual speech recognition.	
• Integrated adapters into the frozen conformer achieving a 4.1% relative WER reduction with zero test-time fine-tuning.	
PageFox   Flask, SQL, React Native, Linear Contextual Bandits, Recommendation System	July 2024
• Designed multi-layered architecture (data models, services, API, client), enabling scalable recommendation serving.	
• Created a relational database using MySQL with SQLAlchemy ORM, enabling efficient data management and retrieval.	
• Built a contextual-bandit recommendation system with modular ranking pipeline; dynamically adjusted recommendations based on user interaction signals, improving click-through proxy metrics.	
RL Adversarial Combat in Minecraft   Deep Q-Learning, Deep Q-Networks, Bootstrapping, Self-Play	March 2024
• Designed a reinforcement learning agent, adapted DQN architectures, and employed bootstrapping and self-play.	
Dermetic   Link   Python, Flask, AWS, Nginx, Scikit-learn, MongoDB, REST	November 2023
• Implemented deep learning skin cancer detection system, hosted using Nginx for reverse proxy, and AWS EC2.	

## Technical Skills

Languages: Python, Java, Scala, JavaScript, HTML, CSS, C/C++, SQL, R	
Technologies/Frameworks: PyTorch, TensorFlow, Numpy, Pandas, Docker, Knowledge Distillation, Graph RAG, Agentic AI, Kubernetes, Flask, MongoDB, Linux, Unix, ReactJS, React Native, AWS, EC2, S3, EFS, Lambda, Spring, REST, Grafana, SciKit Learn, MySQL, SQLite, ONNX, Hugging Face Transformers, Kafka, BigQuery, Accelerate, TensorRT, Optuna, ViT, DETR, Reinforcement Learning, Computer Vision, NLP, Transformers, Edge Computing	